

COM304 COMPUTER ARCHITECTURE

10 January 2019

Dr Noor Mahammad Sk

COM304 Course Evaluation

2

- Pre End Semester: 40% weightage
 - ▣ Quiz 1 – 20% weightage
 - ▣ Quiz 2 – 20% weightage
- End semester: 60% Weightage
 - ▣ End Sem Exam – 60% weightage
- Solution to Open Challenges will be graded 'S'

Syllabus

3

- **Fundamentals of computer design:** Classes of computers, trends in technology, measurement of performance of a computer system, current issues in design of functional components of a computer system – Processor unit, memory unit, and secondary storage unit; Hardware/software tradeoff in computer design
- **Fundamentals of processor design:** Instruction set processor design, exploitation of instruction level parallelism, processor micro architecture, performance of a processor
- **Pipelined processor architecture:** Fundamentals of pipelining, arithmetic pipeline design – Carry look ahead adder, Wallace tree multiplier, Floating–point adder/subtractor;
- **Instruction pipeline design;** Balancing pipeline stages; Stalls in a pipeline; Methods for reductions of stalls in a pipelined processor

Syllabus Contd.,

4

- **Superscalar processor architecture:** Limitations of scalar pipelines, superscalar pipelines, dynamic exploitation of instruction–Level parallelism, register dataflow techniques, memory dataflow techniques, Instruction flow techniques, case studies of superscalar processors
- **Advanced processor architectures:** Multithreaded processors, multi core processors, reconfigurable instruction set processors
- **Storage system architectures:** RAID architecture, storage area networks, Network attached storage
- **Large computer system architectures:** Symmetric multiprocessor systems – Shared memory systems and shared bus architectures; cache coherency protocols – MESI protocol and coherence in multi–level cache systems; Internetwork architectures – Directory protocol for cache coherence

Reference Books

5

- Hennessy J.H and Patterson D.A, Computer Architecture – A Quantitative Approach, Morgan Kaufmann, 2003.
- Shen J.P and Lipasti M.H, Modern Processor Design – Fundamentals of Superscalar Processors, Tata McGraw Hill, 2003.

Class Timings

6

- 100% Attendance must

DAY	Time
Monday	08:00 - 08:50
Thursday	12:00 - 12:50
Friday	09:00 - 09:50

7

Fundamentals of Computer Design


Computer Technology

8

- Performance improvements:
 - ▣ Improvements in semiconductor technology:
 - Feature size, clock speed
 - ▣ Improvements in computer architecture
 - Enabled by HLL compilers, UNIX
 - Lead to RISC architectures
- Together have enabled:
 - ▣ Lightweight computers
 - ▣ Productivity based managed/interpreted programming languages

Single Processor Performance

9



Current Trends in Architecture

10

- ❑ Cannot continue to leverage instruction-level parallelism (ILP)
 - ▣ Single processor performance improvement ended in 2003
- ❑ New models for performance
 - ▣ Data level parallelism (DLP)
 - ▣ Thread-level parallelism (TLP)
 - ▣ Request-level Parallelism (RLP)
- ❑ These require explicit restructuring of the application

Classes of Computers

11

- Personal Mobile Devices (PMD)
 - ▣ e.g., smart phones, tablet computers
 - ▣ Emphasis on energy efficient and real-time
- Desktop Computing
 - ▣ Emphasis on price-performance
- Servers
 - ▣ Emphasis on availability, scalability and throughput
- Clusters/warehouse scale computers
 - ▣ Used for “Software as a Service (SaaS)”
 - ▣ Emphasis on availability and price-performance
 - ▣ Subclass: Super computers
 - Emphasis: floating-point performance and fast internal networks
- Embedded Computers
 - ▣ Emphasis: price

Parallelism

12

- Classes of parallelism in applications:
 - ▣ Data-Level Parallelism (DLP)
 - ▣ Task-Level Parallelism (TLP)

- Classes of architectural parallelism:
 - ▣ Instruction-Level Parallelism (ILP)
 - ▣ Vector architecture/Graphic Processor Units (GPUs)
 - ▣ Thread-Level Parallelism
 - ▣ Request-Level Parallelism

Flynn's Taxonomy

13

- Single instruction stream, single data stream (SISD)
- Single instruction stream, multiple data streams (SIMD)
 - ▣ Vector architectures
 - ▣ Multimedia extensions
 - ▣ Graphics processor units
- Multiple instruction streams, single data stream (MISD)
 - ▣ No commercial implementation
- Multiple instruction streams, multiple data streams (MIMD)
 - ▣ Tightly-coupled MIMD
 - ▣ Loosely-coupled MIMD

Defining Computer Architecture

14

- “Old” view of computer architecture:
 - ▣ Instruction set architecture (ISA) design
 - ▣ i.e., decisions regarding
 - Registers, memory addressing, addressing modes, instruction operands, available operations, control flow instruction, instruction encoding
- “Real” computer architecture:
 - ▣ Specific requirements of the target machine
 - ▣ Design to maximize performance within constraints:
 - Cost, power and availability
 - ▣ Includes ISA, micro-architecture, hardware

Trends in Technology

15

- Integrated circuit technology
 - ▣ Transistor density: 35%/year
 - ▣ Die size: 10-20%/year
 - ▣ Integration overall: 40 – 55%/year
- DRAM capacity: 25–40 % (slowing)
- Flash capacity: 50-60%/year
 - ▣ 15 – 20x cheaper/bit than DRAM
- Magnetic disk technology: 40%/year
 - ▣ 15 – 25x cheaper/bit than Flash
 - ▣ 300 – 500x cheaper/bit than DRAM

Bandwidth and Latency


16

- Bandwidth or throughput
 - ▣ Total work done in a given time
 - ▣ 10,000 – 25,000x improvement for processors
 - ▣ 300 – 1200x improvement for memory and disks

- Latency or response time
 - ▣ Time between start and completion of an event
 - ▣ 30 – 80x improvement for processors
 - ▣ 6 – 8x improvement for memory and disks

Bandwidth and Latency

17



Log-log plot of bandwidth and latency milestones

Transistors and Wires

18

- Feature Size
 - ▣ Minimum size of transistor or wire in x or y dimension
 - ▣ 10 micros in 1971 to 0.032 microns in 2011
 - ▣ Transistor performance scales linearly
 - Wire delay does not improve with feature size!
 - ▣ Integration density scales quadratically

Power and Energy

19

- ❑ Problem: Get power in, Get power out
- ❑ Thermal Design Power (TDP)
 - ▣ Characterizes sustained power consumption
 - ▣ Used as target for power supply and cooling system
 - ▣ Lower than peak power, higher than average power consumption
- ❑ Clock rate can be reduced dynamically to limit power consumption
- ❑ Energy per task is often a better measurement

Dynamic Energy and Power


20

- Dynamic energy
 - ▣ Transistor switch from $0 \rightarrow 1$ or $1 \rightarrow 0$
 - ▣ $\frac{1}{2} \times \text{Capacitive load} \times \text{Voltage}^2$
- Dynamic power
 - ▣ $\frac{1}{2} \times \text{Capacitive load} \times \text{Voltage}^2 \times \text{Frequency switched}$
- Reducing clock rate reduces power, not energy

Power

21

- ❑ Intel 80386 consumed ~ 2 W
- ❑ 3.3 GHz Intel Core i7 consumes 130 W
- ❑ Heat must be dissipated from 1.5×1.5 cm chip
- ❑ This is the limit of what can be cooled by air



Reducing Power

22

- Techniques for reducing power:
 - ▣ Dynamic voltage frequency scaling
 - ▣ Low power state for DRAM, disks
 - ▣ Overclocking, turning off cores

Static Power

23

- Static power consumption
 - ▣ $\text{Current}_{\text{static}} \times \text{Voltage}$
 - ▣ Scales with number of transistors
 - ▣ **To reduce:** power gating

Trends in Cost

24

- Cost driven down by learning curve
 - ▣ Yield

- DRAM: price closely tracks cost

- Microprocessors: price depends on volume
 - ▣ 10% less for each doubling of volume

Integrated Circuit Cost

25

□ Integrated Circuit

$$\text{Cost of Integrated Circuit} = \frac{\text{Cost of die} + \text{Cost of testing die} + \text{Cost of packaging and final test}}{\text{Final test yield}}$$

$$\text{Cost of die} = \frac{\text{Cost of Wafer}}{\text{Die per wafer} \times \text{Die yield}}$$

$$\text{Dies per wafer} = \frac{\pi \times (\text{Wafer diameter}/2)^2}{\text{Die area}} - \frac{\pi \times \text{Wafer diameter}}{\sqrt{2} \times \text{Die area}}$$

□ Bose-Einstein formula:

$$\text{Die yield} = \text{Wafer yield} \times 1 / (1 + \text{Defects per unit area} \times \text{Die area})^N$$

- Defects per unit area = 0.016 – 0.057 defects per square cm (2010)
- N = process-complexity factor = 11.5 – 15.5 (40mm, 2010)

Dependability

26

- Module reliability
 - ▣ Mean time to failure (MTTF)
 - ▣ Mean time to repair (MTTR)
 - ▣ Mean time between failures (MTBF) = $MTTF + MTTR$
 - ▣ Availability = $MTTF / MTBF$

Measuring Performance

27

- Typical performance metrics:
 - ▣ Response time
 - ▣ Throughput
- Speedup of X relative to Y
 - ▣ $\text{Execution time}_Y / \text{Execution time}_X$
- Execution time
 - ▣ Wall clock time: includes all system overheads
 - ▣ CPU time: only computation time
- Benchmarks
 - ▣ Kernels (e.g. matrix multiply)
 - ▣ Toy programs (e.g. sorting)
 - ▣ Synthetic benchmarks (e.g. Dhrystone)
 - ▣ Benchmark suites (e.g. SPEC06fp, TPC-C)

Principles of Computer Design

28

- Take Advantage of Parallelism
 - ▣ e.g. multiple processors, disks, memory banks, pipelining, multiple functional units
- Principle of Locality
 - ▣ Reuse of data and instructions
- Focus on the Common Case
 - ▣ Amdahl's Law

$$\text{Execution time}_{\text{new}} = \text{Execution time}_{\text{old}} \times \left((1 - \text{Fraction}_{\text{enhanced}}) + \frac{\text{Fraction}_{\text{enhanced}}}{\text{Speedup}_{\text{enhanced}}} \right)$$

$$\text{Speedup}_{\text{overall}} = \frac{\text{Execution time}_{\text{old}}}{\text{Execution time}_{\text{new}}} = \frac{1}{(1 - \text{Fraction}_{\text{enhanced}}) + \frac{\text{Fraction}_{\text{enhanced}}}{\text{Speedup}_{\text{enhanced}}}}$$

Principles of Computer Design

29

□ The Processor Performance Equation

$$\text{CPU time} = \text{CPU clock cycles for a program} \times \text{Clock cycle time}$$

$$\text{CPU time} = \frac{\text{CPU clock cycles for a program}}{\text{Clock rate}}$$

$$\text{CPI} = \frac{\text{CPU clock cycles for a program}}{\text{Instruction count}}$$

$$\text{CPU time} = \text{Instruction count} \times \text{Cycles per instruction} \times \text{Clock cycle time}$$

$$\frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock cycle}} = \frac{\text{Seconds}}{\text{Program}} = \text{CPU time}$$

Principles of Computer Design

30

- Different instruction types having different CPIs

$$\text{CPU clock cycles} = \sum_{i=1}^n \text{IC}_i \times \text{CPI}_i$$

$$\text{CPU time} = \left(\sum_{i=1}^n \text{IC}_i \times \text{CPI}_i \right) \times \text{Clock cycle time}$$