


Complexity of PR – An Example

Problem: Sorting incoming fish on a conveyor belt.

Assumption: Two kind of fish:

- (1) sea bass
- (2) salmon



salmon



sea bass



salmon



salmon



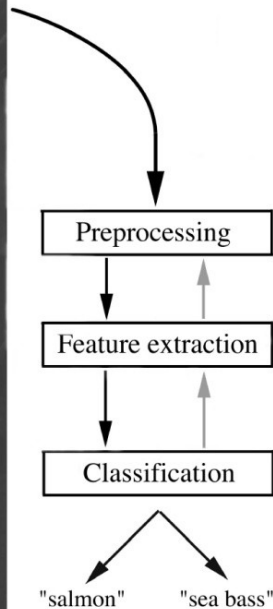
sea bass



sea bass



Pre-processing Step



Example

(1) Image enhancement


(2) Separate touching or occluding fish

(3) Find the boundary of each fish

Feature Extraction

- Assume a fisherman told us that a sea bass is generally **longer** than a salmon.
- We can use **length** as a feature and decide between sea bass and salmon according to a **threshold** on length.
- **How** should we choose the threshold?

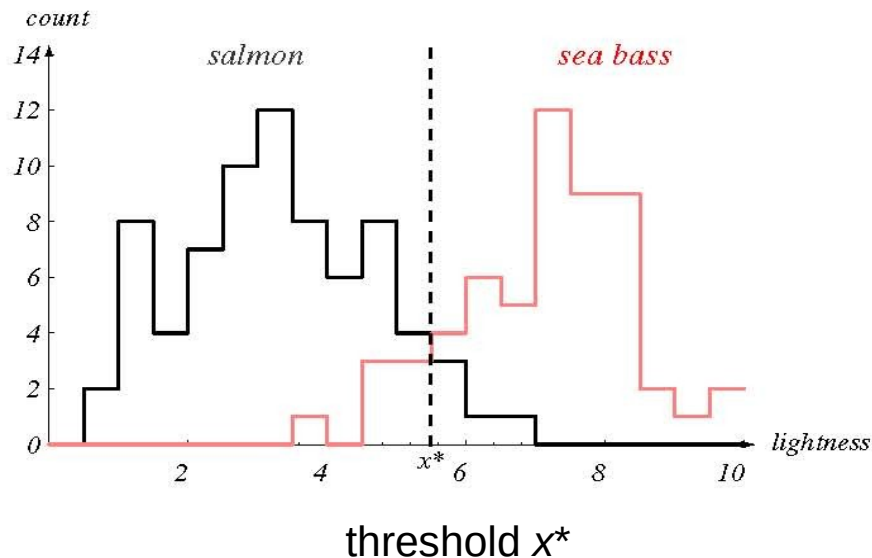
“Length” Histograms



- Even though sea bass is longer than salmon on the average, there are many examples of fish where this observation does not hold.

“Average Lightness” Histograms

- Consider a different feature such as “average lightness”



- It seems easier to choose the threshold x^* but we still cannot make a perfect decision.

Multiple Features

- To improve recognition accuracy, we might have to use more than one features at a time.
 - Single features might not yield the best performance.
 - Using combinations of features might yield better performance.

$$\begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

- x_1 : length
- x_2 : lightness

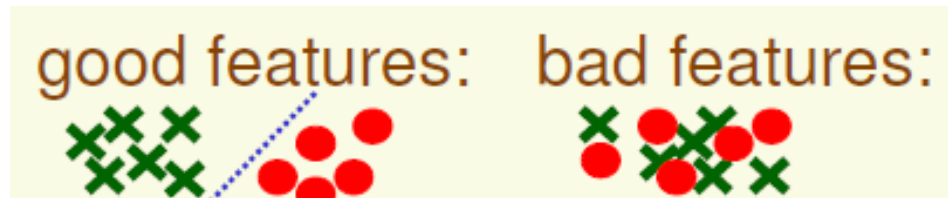
- **How** many features should we choose?

How Many Features?

- Does adding more features always improve performance?
 - It might be **difficult** and **computationally expensive** to extract certain features.
 - **Correlated** features might not improve performance.
 - **“Curse”** of dimensionality.

Feature Extraction


- How to choose a good set of features?
 - Discriminative features



- Invariant features (e.g., translation, rotation and scale)
- Are there ways to automatically learn which features are best ?

Curse of Dimensionality


- Adding **too many** features can, paradoxically, lead to a **worsening** of performance.
 - Divide each of the input features into a number of intervals, so that the value of a feature can be specified approximately by saying in which interval it lies.



- If each input feature is divided into M divisions, then the total number of cells is M^d (d : # of features).
 - Since each cell must contain at least one point, the number of training data grows **exponentially** with d .


Curse of Dimensionality- definition

- As the number of features or dimensions grows, **the amount of data we need to generalize accurately grows exponentially."**
- In applied maths, COD refers to the **problem caused by the exponential increase in volume associated with adding extra dimensions to a mathematical space.**



Curse of Dimensionality-(contd)

- Fig. 1 (a) shows 10 data points in one dimension i.e. there is only one feature in the data set. It can be easily represented on a line with only 10 values, $x=1, 2, 3 \dots 10$.



Curse of Dimensionality-(contd)


- But if we add one more feature, same data will be represented in 2 dimensions (Fig.1 (b)) causing increase in dimension space to $10 * 10 = 100$.
- And again if we add 3rd feature, dimension space will increase to $10 * 10 * 10 = 1000$. As dimensions grows, dimensions space increases exponentially.
- $10^1 = 10$
- $10^2 = 100$
- $10^3 = 1000$ and so on...

Missing Features

- Certain features might be missing (e.g., due to occlusion).
- How should we train the classifier with missing features ?
- How should the classifier make the best decision with missing features ?

Classification

- Partition the *feature space* into two regions by finding the **decision boundary** that minimizes the error.



- How should we choose the decision boundary?

Main Classification Approaches

x : input vector (pattern)



y : class label (class)

- *Generative*


- Model the joint probability, $p(x, y)$
- Make predictions by using Bayes rules to calculate $p(y|x)$
- Pick the most likely label y

- *Discriminative*

- Estimate $p(y|x)$ directly (e.g., learn a direct map from inputs x to the class labels y)
- Pick the most likely label y

Complexity

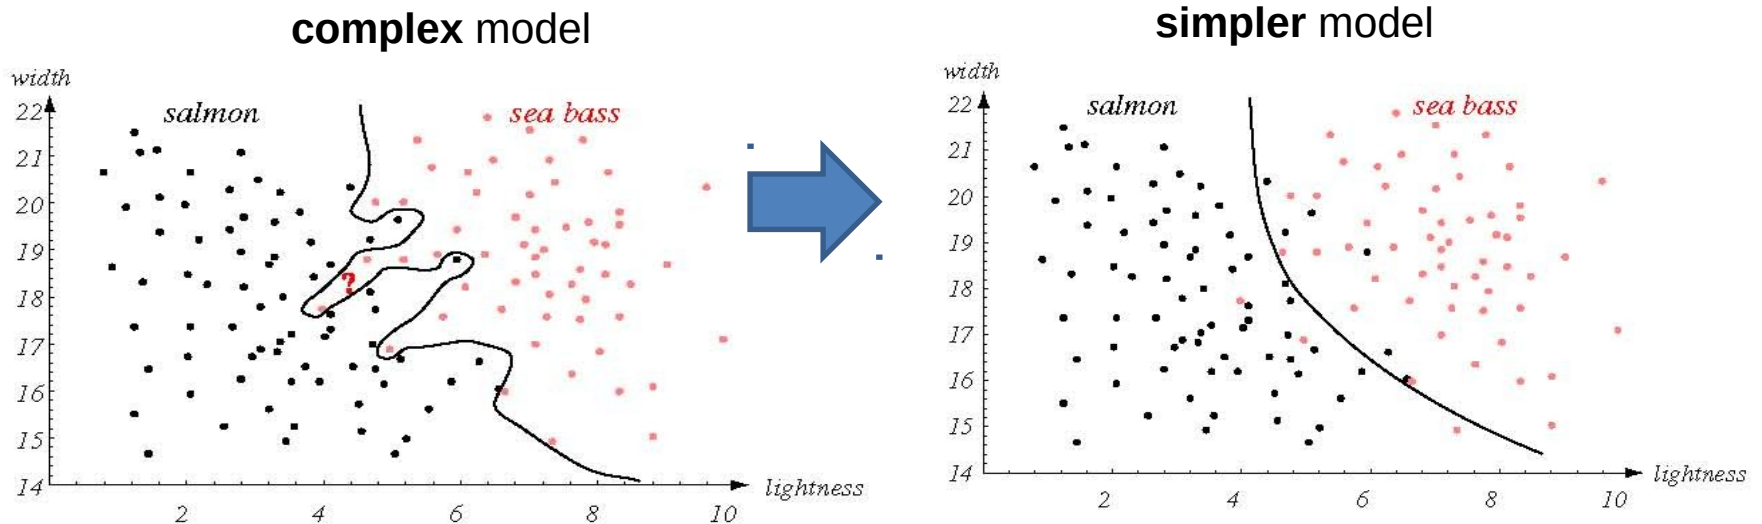
- We can get perfect classification performance on the training data by choosing **complex models**.
- Complex models are **tuned** to the particular training samples, rather than on the characteristics of the true model.



How well can the model **generalize** to unknown samples?

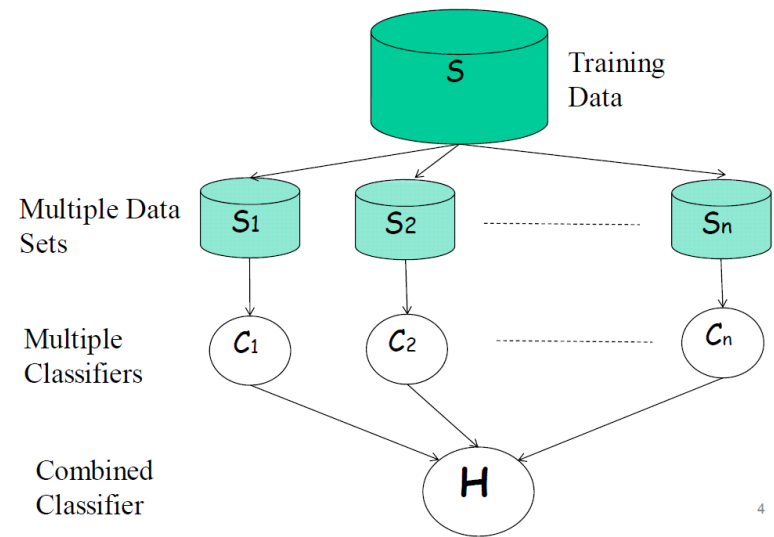
Generalization

- Generalization is defined as the ability of a classifier to produce correct results on **novel** patterns.
- How can we improve generalization performance ?
 - **More** training examples (i.e., better model estimates).
 - **Simpler** models usually yield better performance.



Ensembles of Classifiers

- Performance can be improved using a "pool" of classifiers.
- How should we **build** and **combine** different classifiers ?



Would it be possible to build a “general purpose” PR system?

- Humans have the ability to switch rapidly and seamlessly between different pattern recognition tasks.
- It is very **difficult** to design a system that is capable of performing a variety of classification tasks.
 - Different decision tasks may require different features.
 - Different features might yield different solutions.
 - Different tradeoffs exist for different tasks.